

# Lecture 4: Containers II

AC215

Pavlos Protopapas  
SEAS/Harvard



# Outline

---

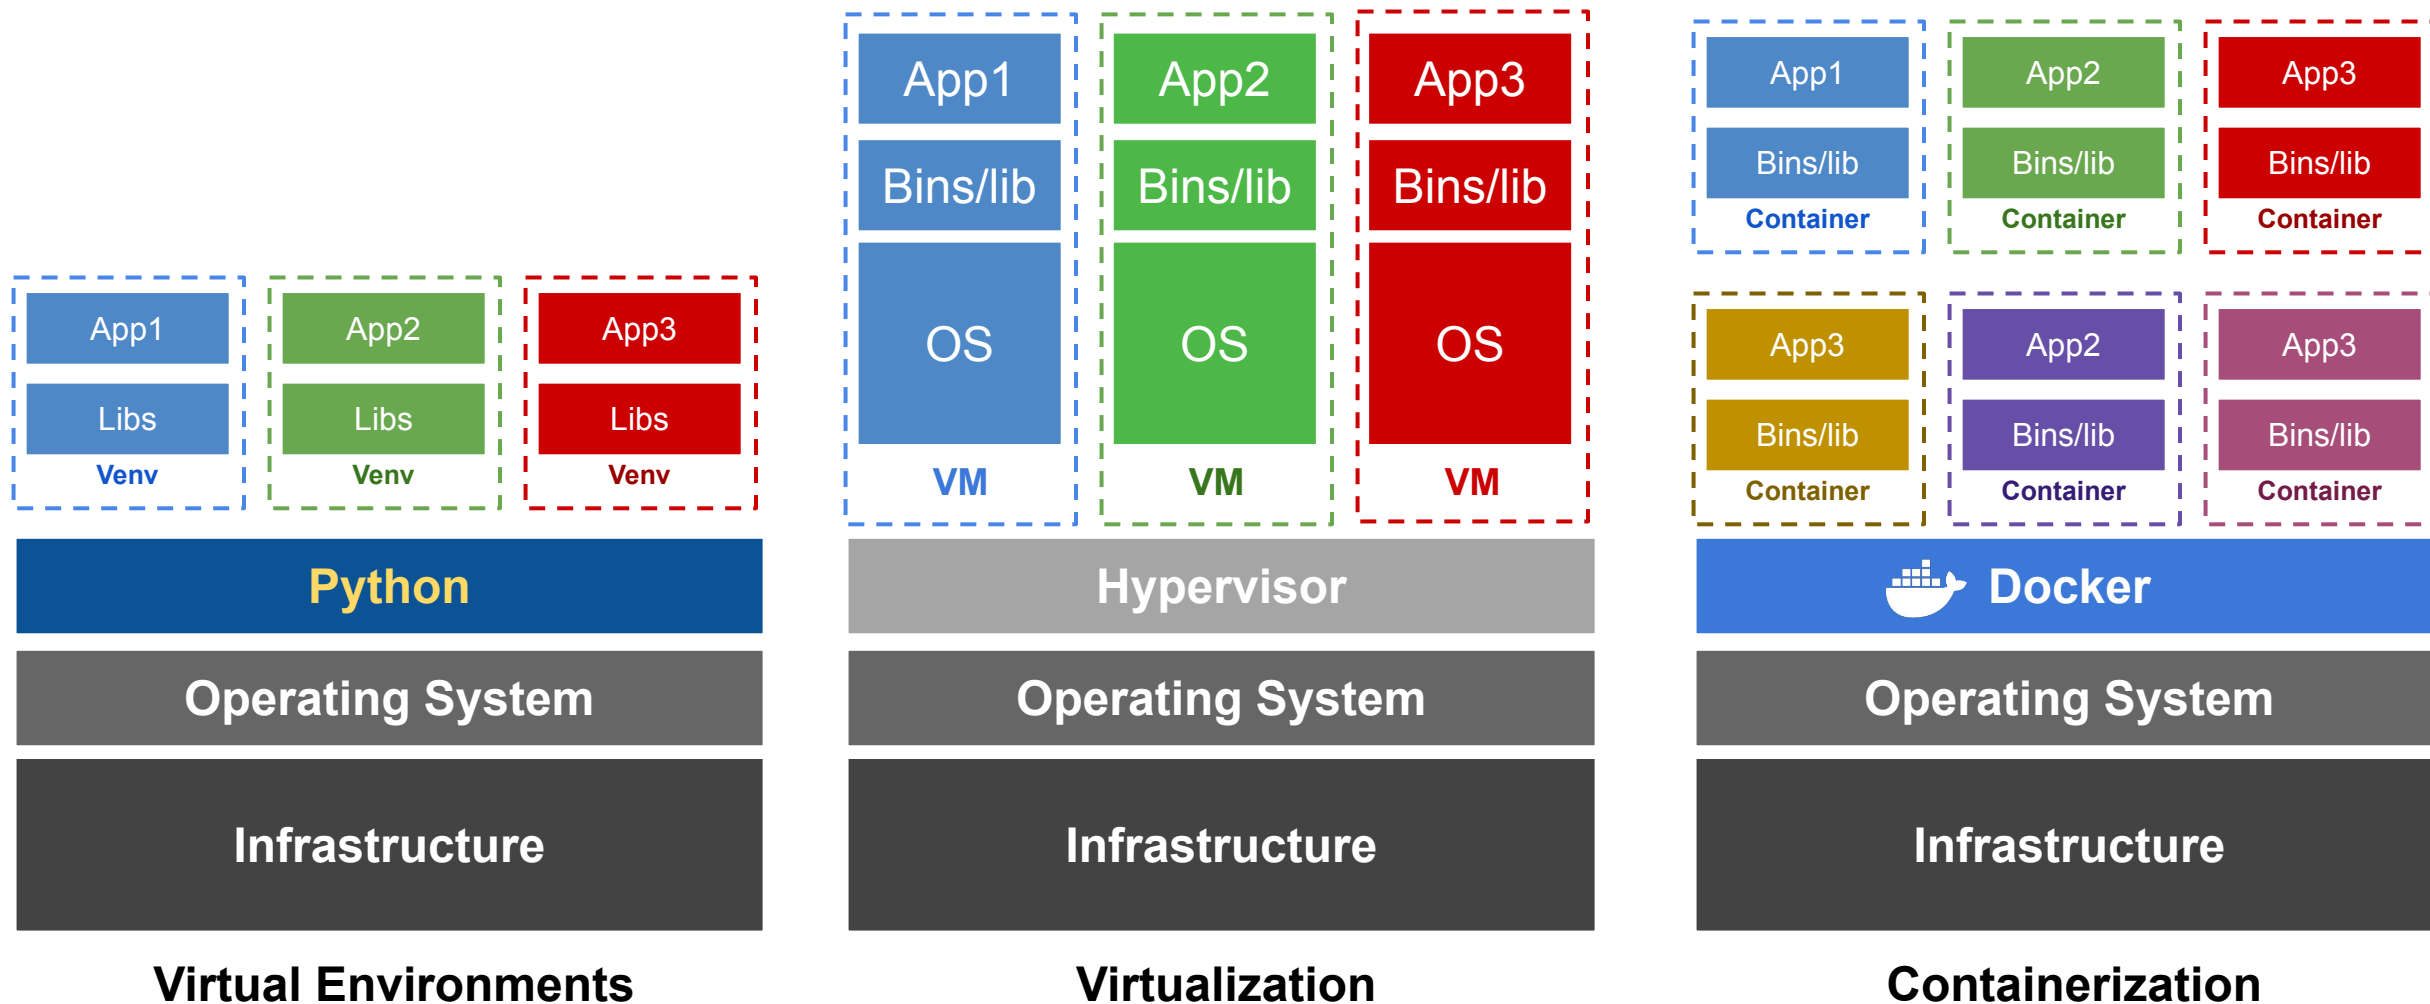
1. Recap: Review of Previous Material
2. Containers in Architecture: Microservices vs. Monolithic
3. Implementing Containers as Microservices
4. Working with Containers Workflow
5. Containers: Pro-tips

# Outline


---

1. **Recap: Review of Previous Material**
2. Containers in Architecture: Microservices vs. Monolithic
3. Implementing Containers as Microservices
4. Working with Containers Workflow
5. Containers: Pro-tips

# Recap: Environments vs Virtualization vs Containerization



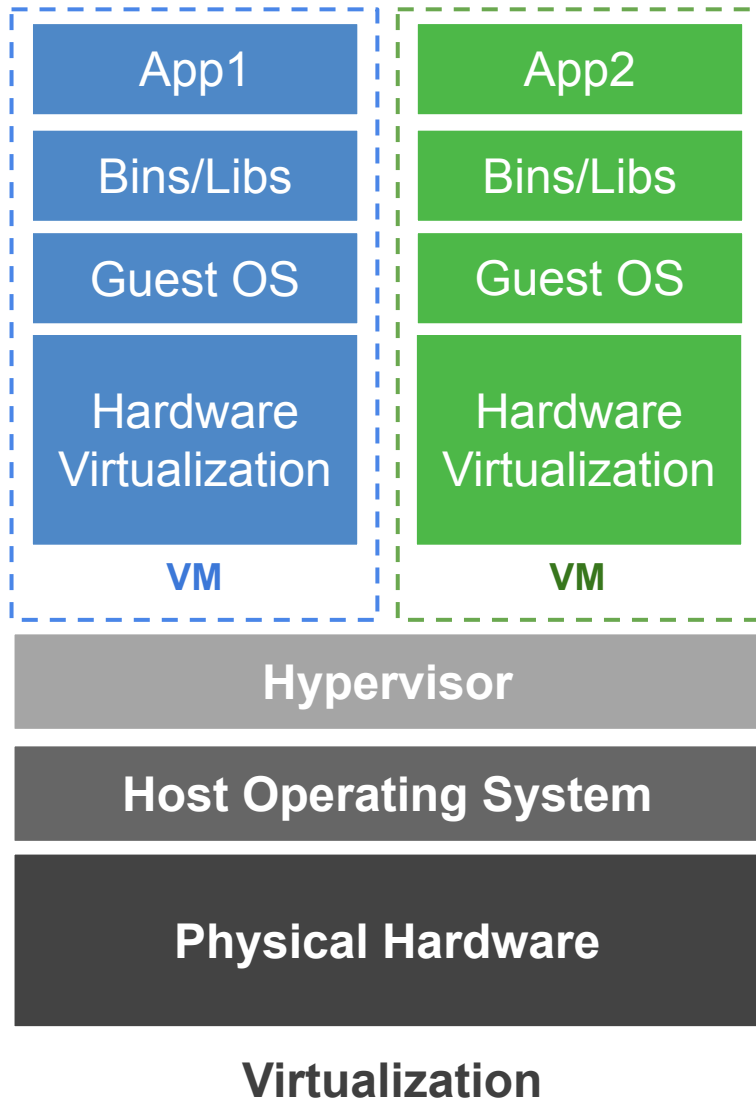
# Environments



Virtual Environments

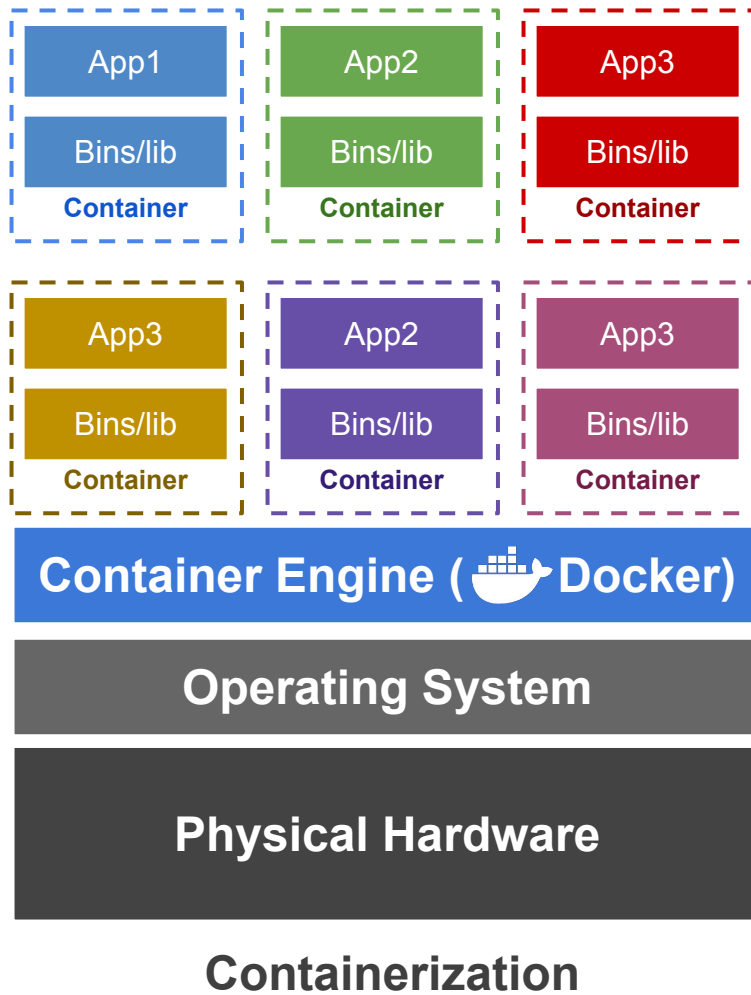
- **Dependency Isolation:** Virtual environments redirect dependencies to their own directories, avoiding system-wide installs.
- **No Kernel Isolation:** Unlike VMs or containers, they don't isolate the kernel.
- **Resource Efficiency:** Without an OS or kernel, virtual environments are lightweight and resource-efficient.
- **Filesystem Access:** Files written within a virtual environment can be accessed from other environments, as there's no filesystem isolation.

# Virtualization (Virtual Machines)



- **CPU Virtualization:** Virtual CPUs are mapped to physical cores, but hypervisor management adds some overhead.
- **Emulated Devices:** VMs use virtual devices (CPUs, network adapters, disks) translated by the hypervisor to real hardware.
- **Full OS:** Each VM runs its own guest OS with independent kernel and user spaces, but this reduces efficiency.
- **Resource Allocation:** RAM, CPU, and disk space are often allocated in fixed blocks, limiting flexibility in resource usage.

# Containerization



- **Namespaces:** Containers isolate processes and resources, making them act like independent systems. For example, PID namespaces separate process IDs, and mount namespaces provide unique file systems.
- **Cgroups:** These limit CPU, memory, and IO usage for each container, ensuring efficient resource use.
- **Process Virtualization:** Namespaces and cgroups work together to isolate and control processes.
- **Shared Kernel:** Containers use the host's OS kernel but have their own files, making them lightweight and efficient.
- **Direct Access:** Containers interact with host resources directly, reducing overhead compared to VMs.

# Outline

---

1. Recap: Review of Previous Material
2. **Containers in Architecture: Microservices vs. Monolithic**
3. Implementing Containers as Microservices
4. Working with Containers Workflow
5. Containers: Pro-tips



# Why use Containers?

---

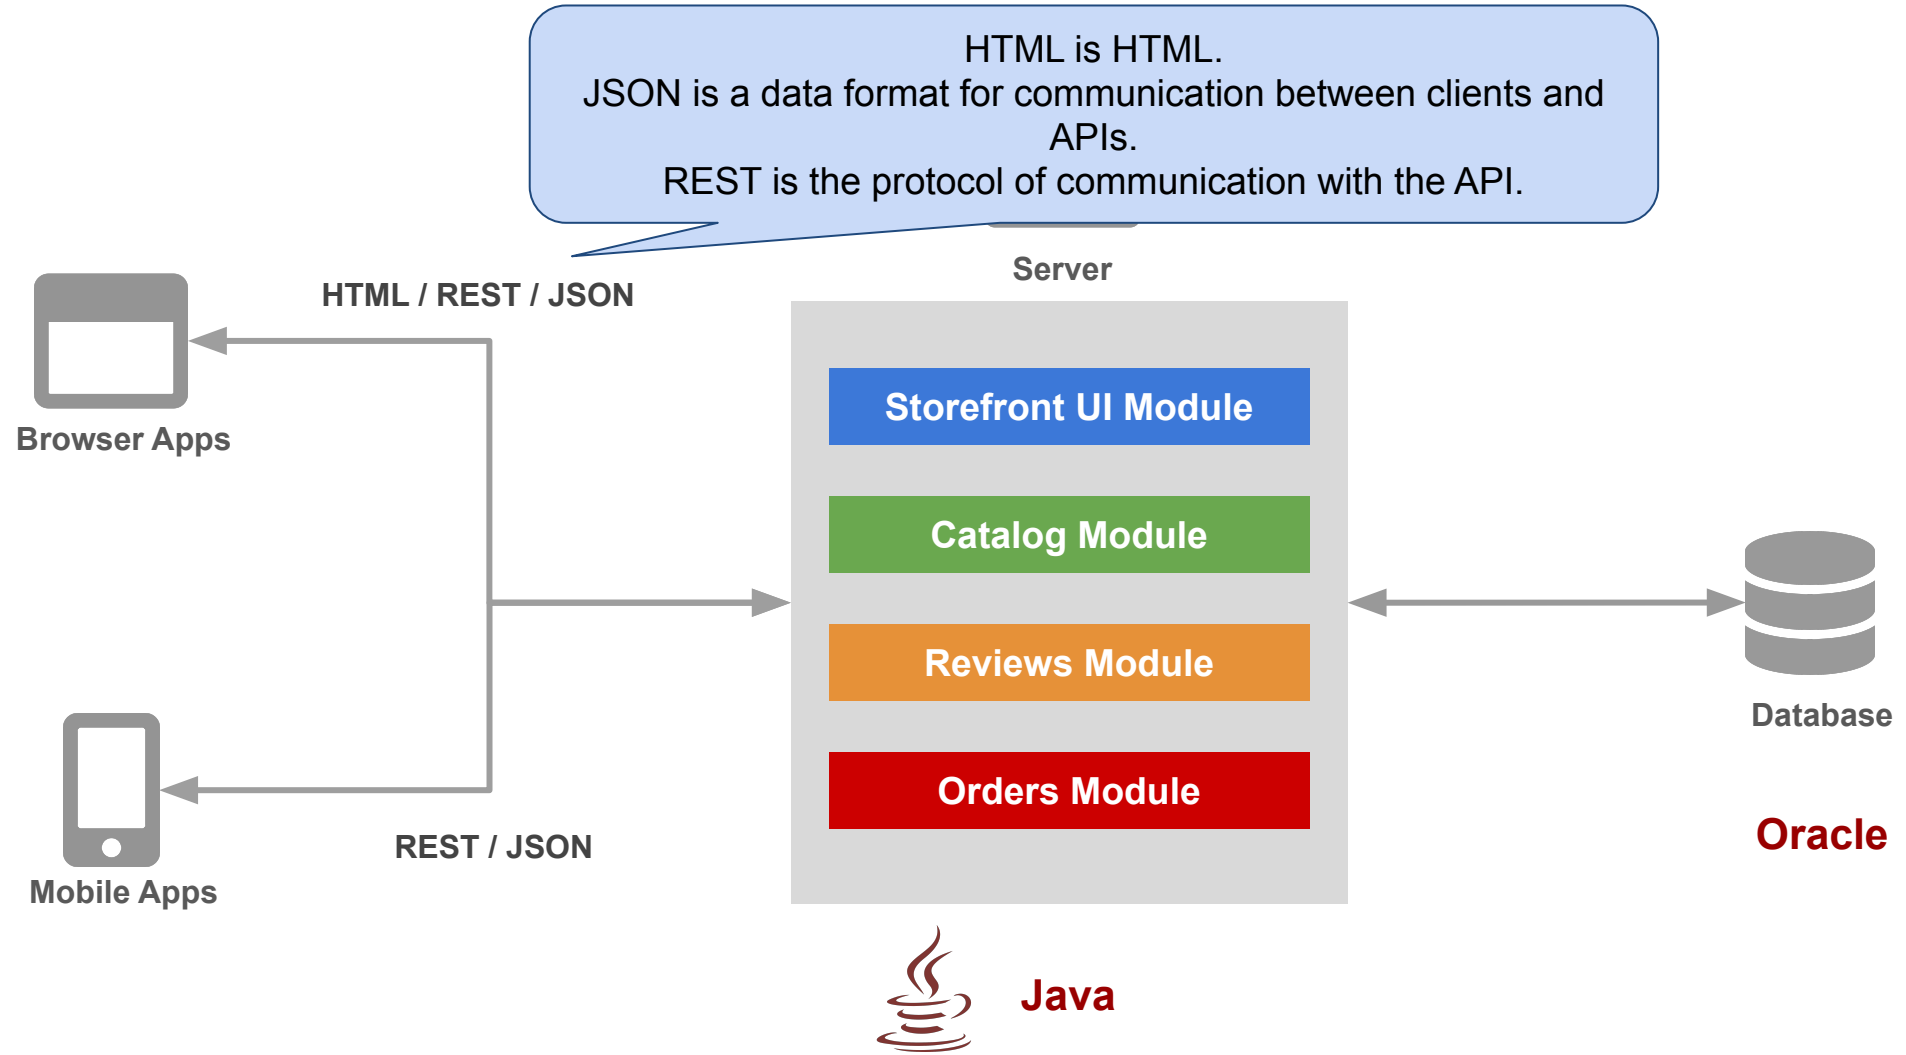
## Conceptual Scenario

- Picture building an application, such as an online cheese store.

## Traditional Approach

- Traditionality you would build this using a [Monolithic Architecture](#).

# Monolithic Architecture



# Monolithic Architecture - **Advantages**

---

## **Simplicity in Development:**

Most tools and IDEs natively support monolithic applications.

## **Ease of Deployment:**

All components bundled into a single, unified package.

## **Scalability:**

Easier to scale by replicating the entire application as a whole (horizontal scaling).

# Monolithic Architecture - Disadvantages

---

## **Maintenance Challenges:**

Complexity increases over time, making it harder to implement changes or find issues.

## **System Vulnerability:**

A failure in a single component can lead to the collapse of the entire system.

## **Patching Difficulties:**

Patching or updating specific modules can be cumbersome due to tightly-coupled components.

# Monolithic Architecture - Disadvantages

---

## Technology Lock-in:

Adopting new technologies or updating existing ones can be problematic due to interdependencies.

## Slow Startup:

Increased startup time as all components must be initialized simultaneously.

## Onboarding Challenges:

New users need to familiarize themselves with the entire codebase.

# Applications have changed dramatically

---

## **A decade ago**

Apps were monolithic  
Built on a single stack (e.g. .Net or Java)  
Long lived  
Deployed to a single server


## **Today**

Apps are constantly being developed  
Build from loosely coupled components  
Newer version are deployed often  
Deployed to a multitude of servers

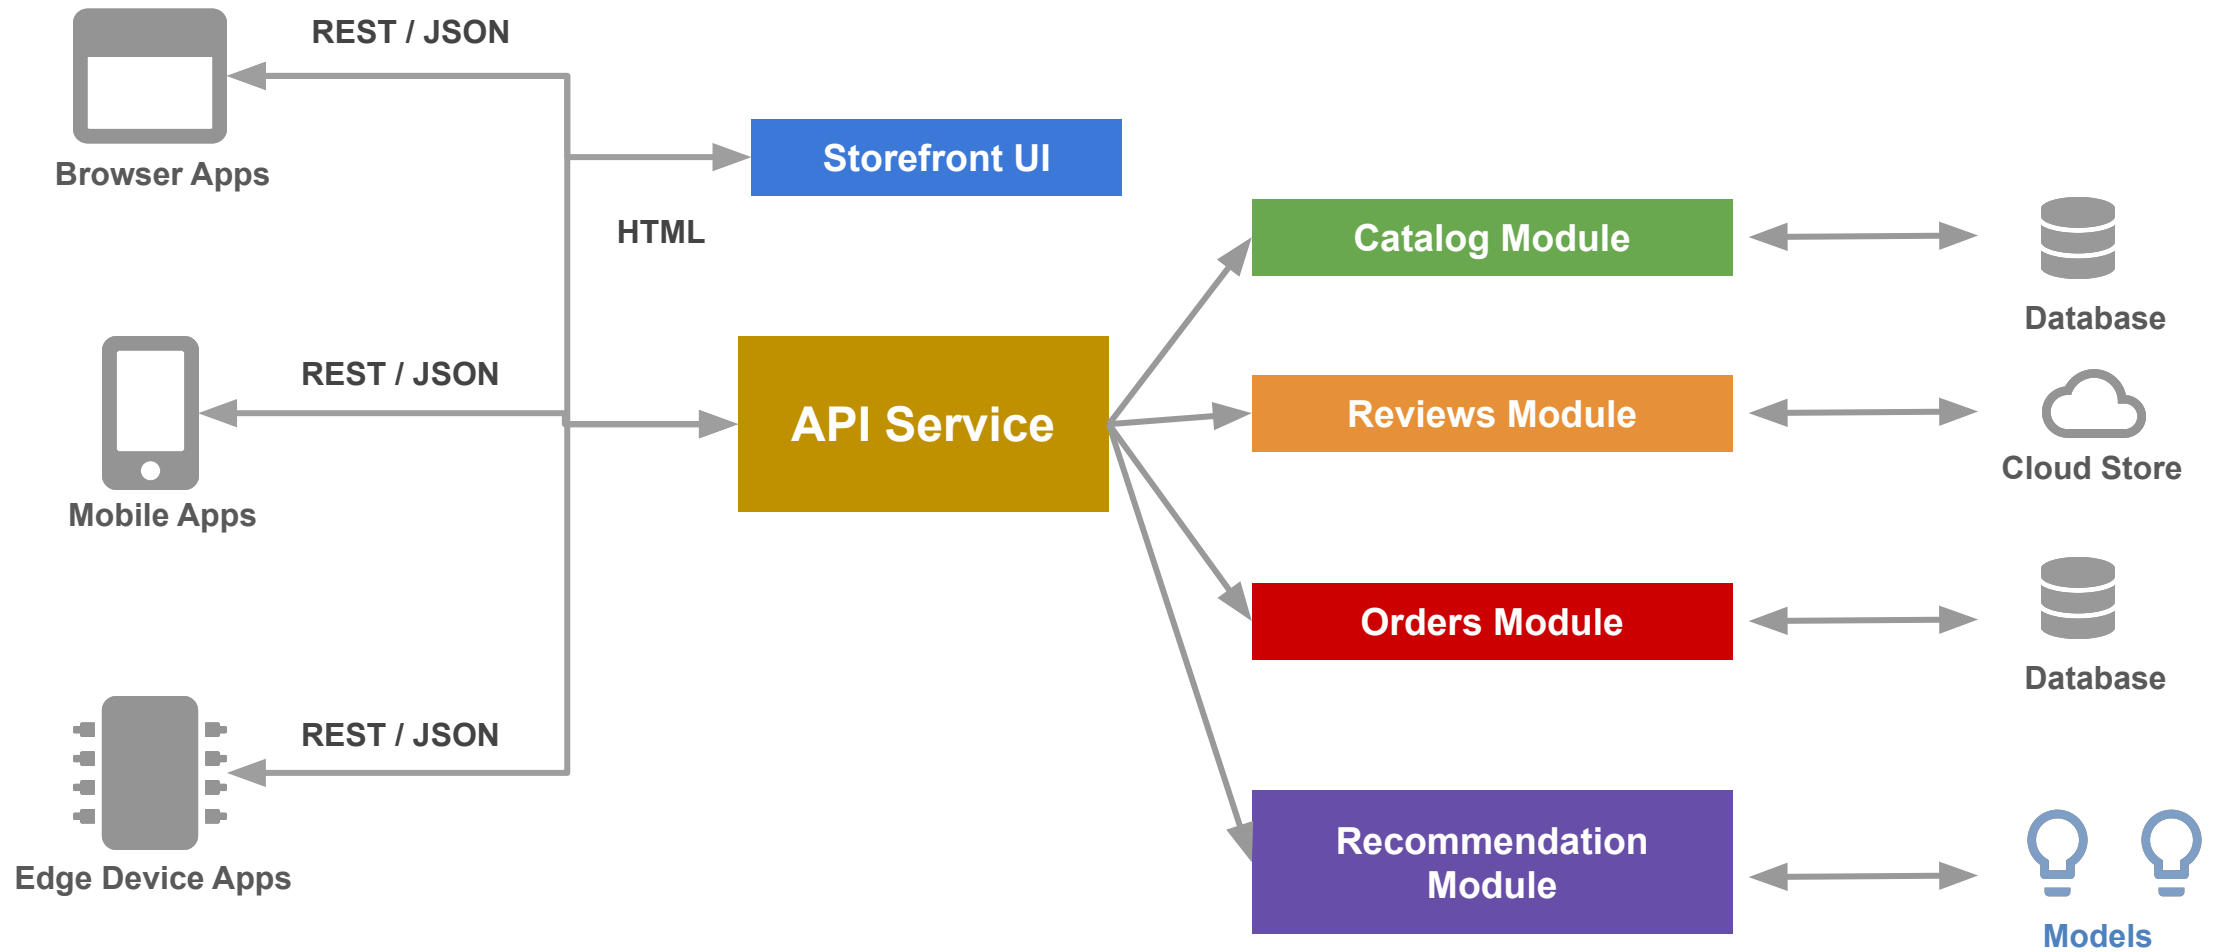
## **Data Science**

**Apps are being integrated with various  
data types/sources and models**

# Monolithic Architecture



# Today: Microservice Architecture





# Microservice Architecture - **Advantages**

---

## **Simplified Maintenance:**

Modular design makes it easier to manage, update, and debug individual services.

## **Fault Isolation:**

Independent components ensure that failure in one service doesn't bring down the entire application.

## **Streamlined Patching:**

Easier to patch or update specific services without affecting the entire system.

## **Technological Flexibility:**

Adapting to or adopting new technologies becomes seamless due to service independence.

## **Quick Startup:**

Reduced startup time as all components can be initialized in parallel.

# Microservice Architecture - Disadvantages

---

## Development and Deployment Complexity:

Using multiple technologies across components can complicate both development and deployment, as managing diverse dependencies requires a more intricate setup.

## Scaling Concerns:

Scaling the entire application can be intricate due to disparate components.


**Docker + Kubernetes**

# Why use Containers?


---

- Consider a software development team workflow for developing an App
- Traditionality you would develop/build this independently in various machines (dev, test, qa, prod)

# Software Development Workflow (no Docker)



# Software Development Workflow (with Docker)



**Containers** need to be setup only once.

# Software Development Workflow (with Docker)



Who creates the Dockerfile, and where is it stored? Do we use pre-built images or does each developer build them? Who is in charge of managing this? Also, what's the process for handling the Pipfile and Pipfile.lock?

Development machines only needs **Docker installed**.

**Containers** need to be setup only once.

This seems like a lot.

# Software Development Workflow (with Docker)




Development machines only needs **Docker installed**.

**Containers** need to be setup only once.




So

Every  
code



Production/ Test Servers



Production server only needs **Docker installed**.

Production server pulls Docker **images** from **container registry** and runs them.

# Outline

---

1. Recap: Review of Previous Material
2. Containers in Architecture: Microservices vs. Monolithic
- 3. Implementing Containers as Microservices**
4. Working with Containers Workflow
5. Containers: Pro-tips



# Tutorial (T5) - Building the Mega Pipeline App

---


One of [Formaggio.me](https://formaggio.me)'s goals is to create a podcast on various cheese-related topics. After recording the podcast, we plan to transcribe the audio, use a language model to correct grammar and enhance the text, and then generate audio that will be made available to our users. Remember, we aim to reach an audience all over the world, so the podcast will be translated into various languages and synthesized into audio.

The goal here is to simulate a realistic development scenario where each component will be developed by different teams and containerized.

**BONUS:** You can use elevenlabs API to generate text with Pavlos' voice or your own voice.



# Tutorial (T5) - Building the Mega Pipeline App



# Tutorial (T5) - Team Challenge!


---

- We'll form teams of 5, and each team must complete 5 tasks.
- The first team to finish all tasks will win a special prize!
- Don't worry, every team that completes the tasks will also get a reward!
- The rewards are a surprise—so give it your best effort and have fun!




# Tutorial (T5) - Building the Mega Pipeline App


≡ AC215: Mega Pipeline App




Audio Prompts




Transcribed Audio




Generated Text




Synthesised Audio



Translated Text



Synthesised Audio



▶ 0:00 — 🔊

**Group: group-01**

welcome to the very cheesy  
podcast where we  
discuss all things cheese this is  
your house  
spotless protopapas we all love  
c... [show more](#)

**Group: staff**

welcome to the very cheesy  
podcast where we  
discuss all things cheese this is  
your house  
spotless protopapas we all love  
c... [show more](#)

**Group: group-01**

Welcome to the Very Cheesy Podcast,  
where  
we discuss all things cheese. This is  
your host,  
Pavlos Protopapas. We all love c...  
[show more](#)

**Group: staff**

Welcome to the Very Cheesy Podcast,  
where  
we discuss all things cheese. This is  
your host,  
Pavlos Protopapas. We all love c...  
[show more](#)

**Group:group-01**

▶ 0:00 — 🔊

**Group:staff**

▶ 0:00 — 🔊

**Group: group-01**

Bienvenue dans le podcast très ringard,  
où  
nous discutons de tout ce qui concerne  
le  
fromage.Ceci est votre hôte, Pav...  
[show more](#)

**Group: staff**

Bienvenue dans le podcast très ringard,  
où  
nous discutons de tout ce qui concerne  
le  
fromage.Ceci est votre hôte, Pav...  
[show more](#)

**Group:group-01**

▶ 0:00 — 🔊

**Group:staff**

▶ 0:00 — 🔊

**Group:group-01**

▶ 0:00 — 🔊

**Group:staff**

▶ 0:00 — 🔊

# GCP Authentication Methods

---

Before we start let us review how do we authenticate to different services/accounts and APIs

## **OAuth 2.0:**

For user-driven authentication and access.

## **Service Account (part of IAM - Identity and Access Management- in GCP):**

For server-to-server interactions requiring automation and high control without user intervention.

## **API Key:**

For lightweight, less secure access to APIs, use cautiously.

## **Default Service Accounts (part of IAM in GCP):**

For Compute Engine, Kubernetes Engine, and App Engine with predefined permissions.

## **Workload Identity Federation:**

For external identities to access GCP securely.

# Tutorial (T5) - Building the Mega Pipeline App

- App: <https://ac215-mega-pipeline.dlops.io/>
- Teams
  - 📝 Task A [transcribe\\_audio](#):
  - 📄 Task B [generate\\_text](#):
  - 🎧 Task C [synthesis\\_audio\\_en](#):
  - 🇫🇷 Task D [translate\\_text](#):
  - 🎧 Task E [synthesis\\_audio](#):
- Instructions: <https://github.com/dlops-io/mega-pipeline>

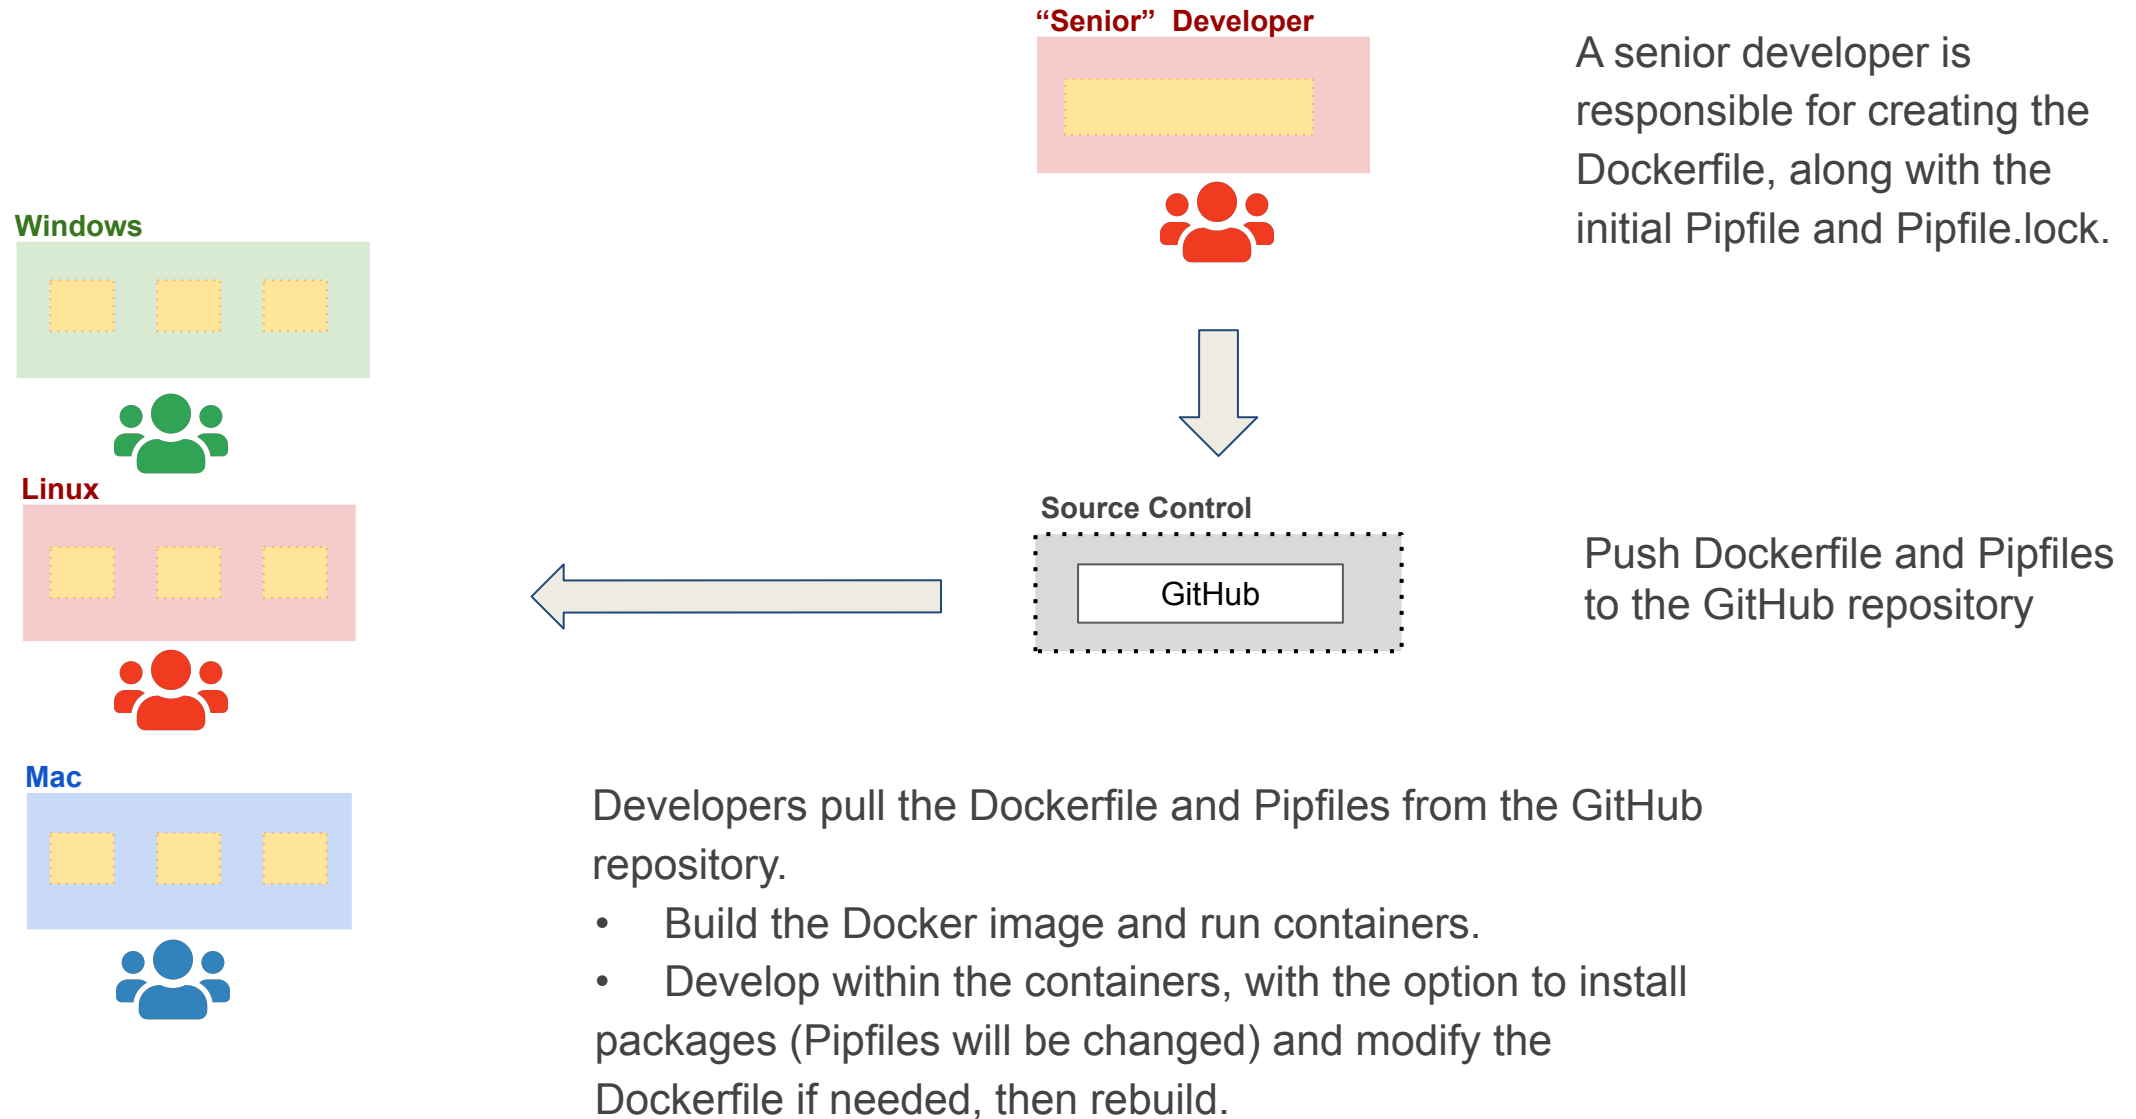


# Outline

---

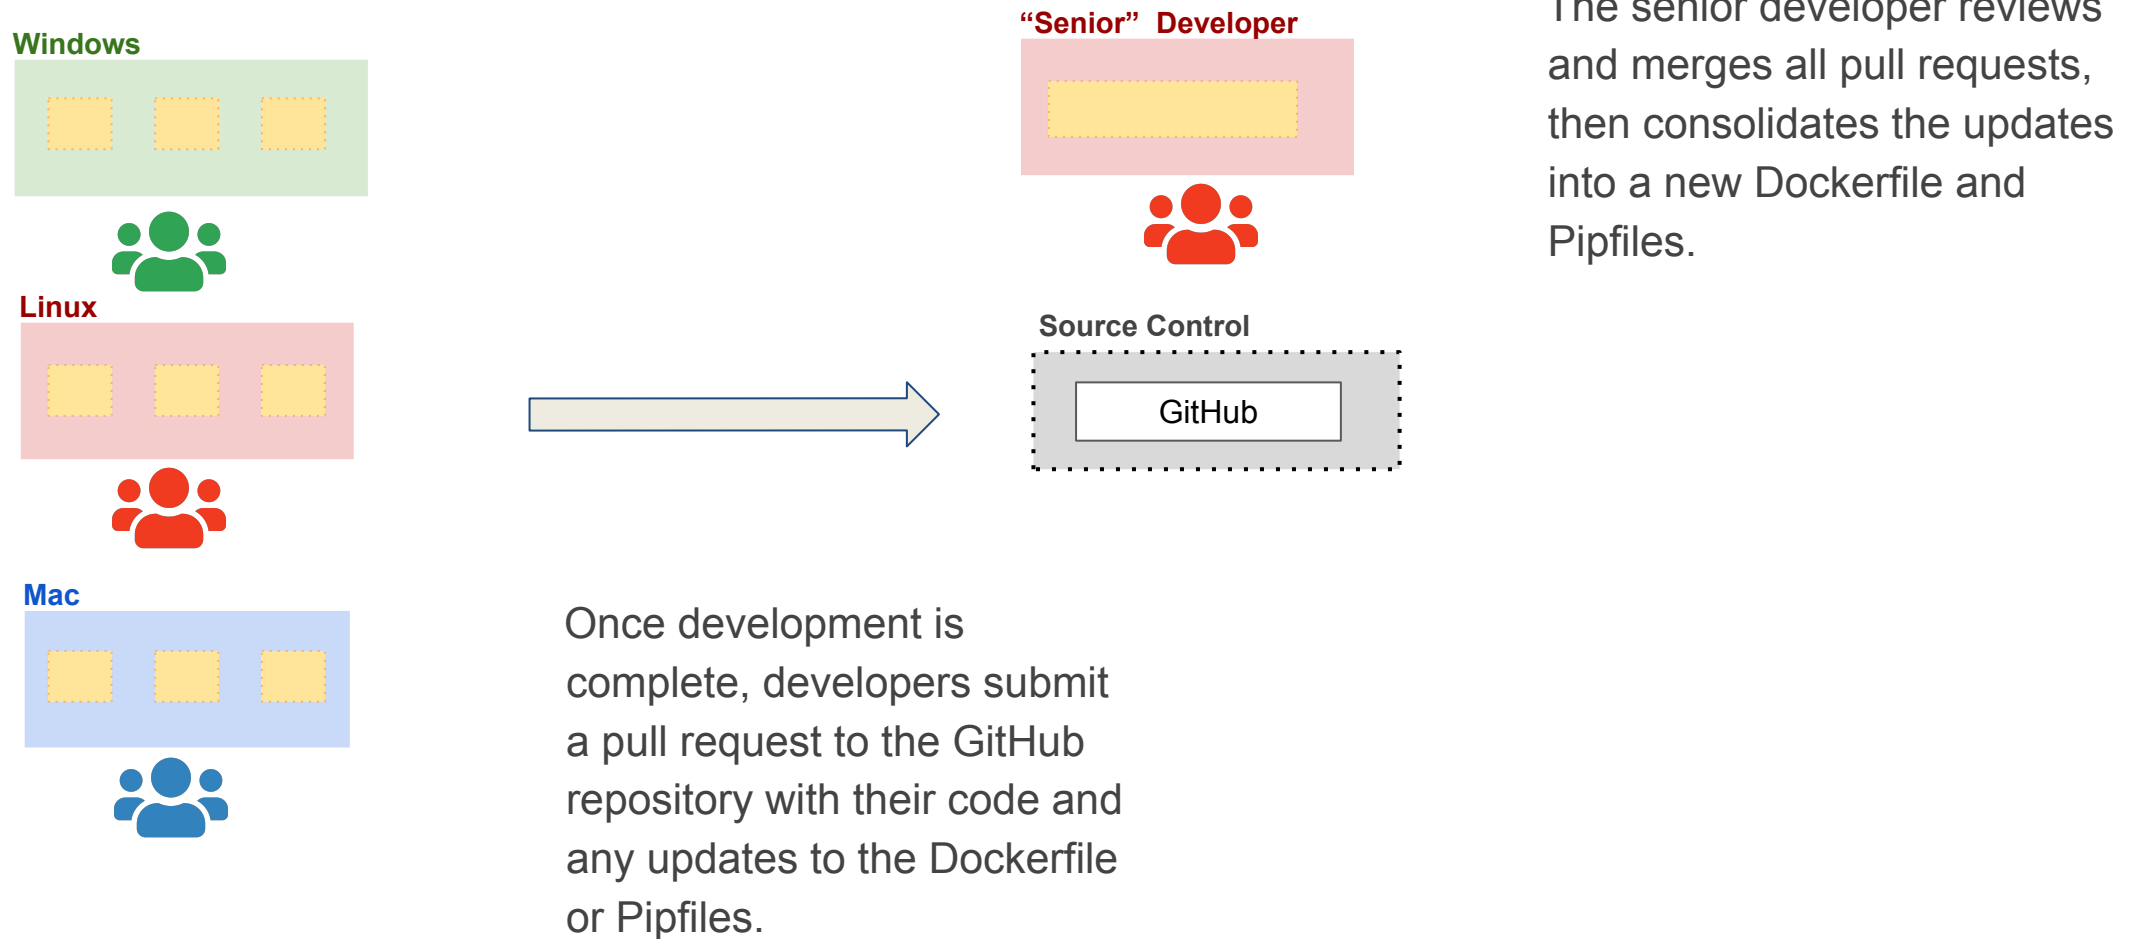
1. Recap: Review of Previous Material
2. Containers in Architecture: Microservices vs. Monolithic
3. Implementing Containers as Microservices
4. **Working with Containers Workflow**
5. Containers: Pro-tips

# Workflow with Docker: **Scenario 1** (early stages of development)

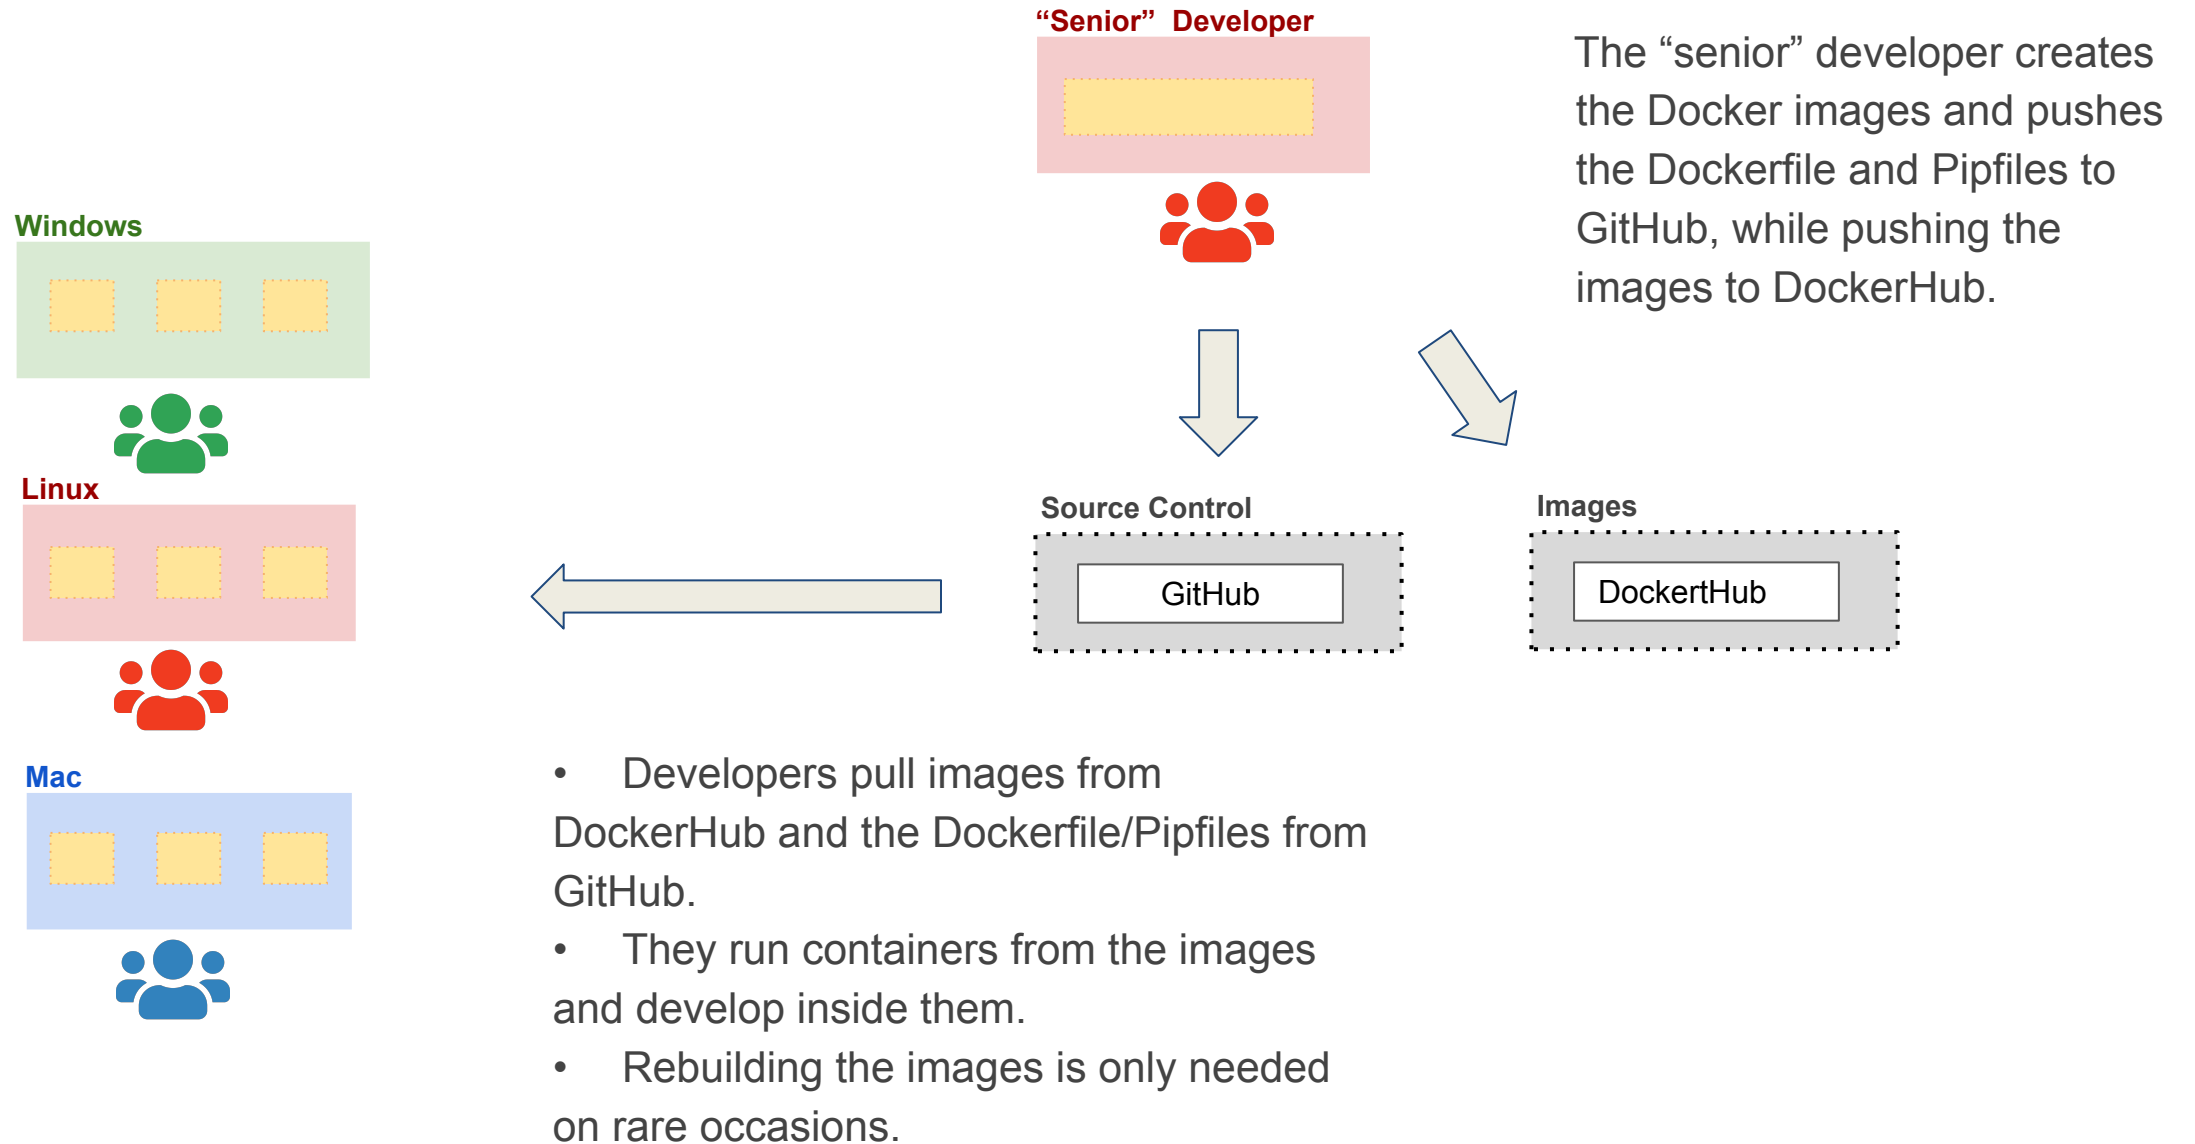




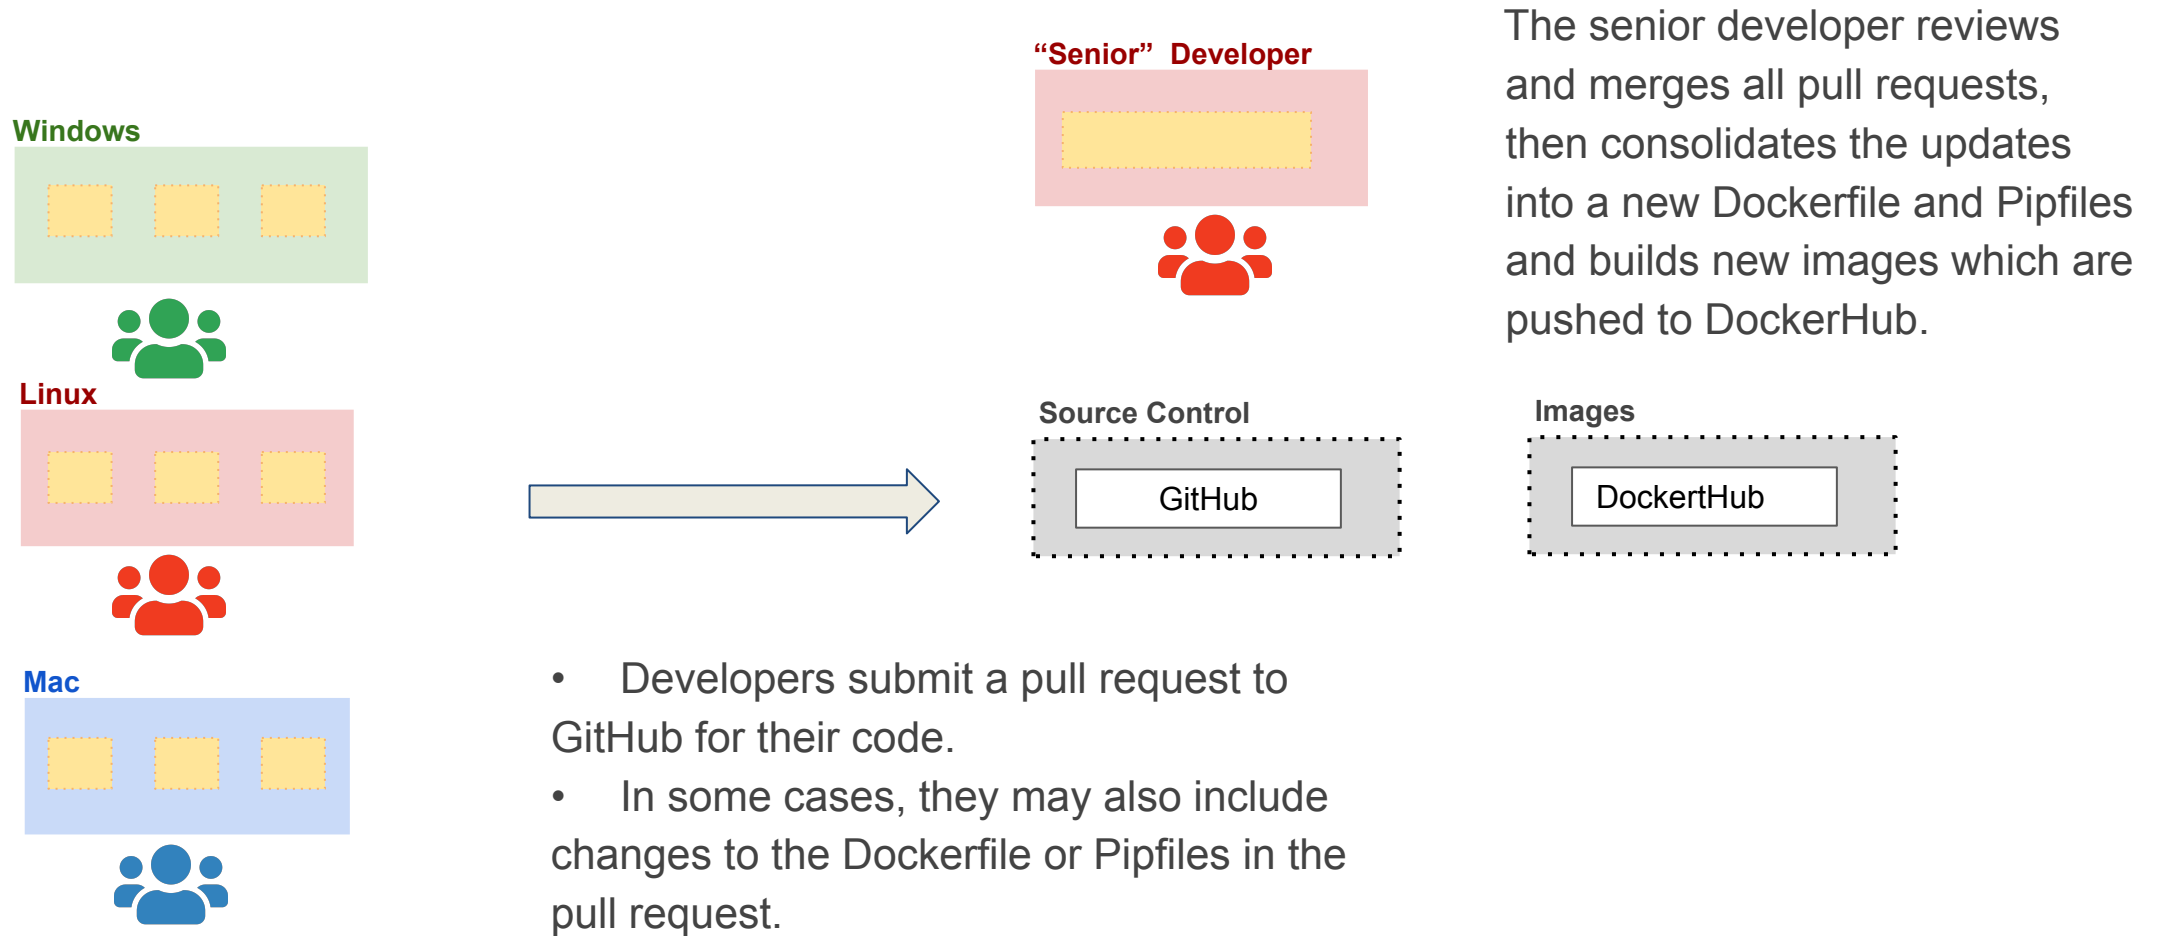
# Workflow with Docker: **Scenario 1** (early stages of development)



# Workflow with Docker: **Scenario 2** (later stages of development)



# Workflow with Docker: **Scenario 2** (later stages of development)



# Workflow with Docker: **A Flexible Approach**

---

**Is there a “prefect” workflow?**

No

**So, how do we decide what to do?**

Clear communication and rules are essential. Each team can have its own workflow, based on the project and team needs.

# Outline

---

1. Recap: Review of Previous Material
2. Containers in Architecture: Microservices vs. Monolithic
3. Implementing Containers as Microservices
4. Working with Containers Workflow
5. **Containers: Pro-tips**

# Pro Tip 1: multi-stage builds

---

Running commands during the build process can take significant disk space. For example, when **installing** and **building** packages, we **download** and generate, numerous files, which can substantially **increase** the **size** of the Docker image.

**Question:** Do we really need all of these files for our app?

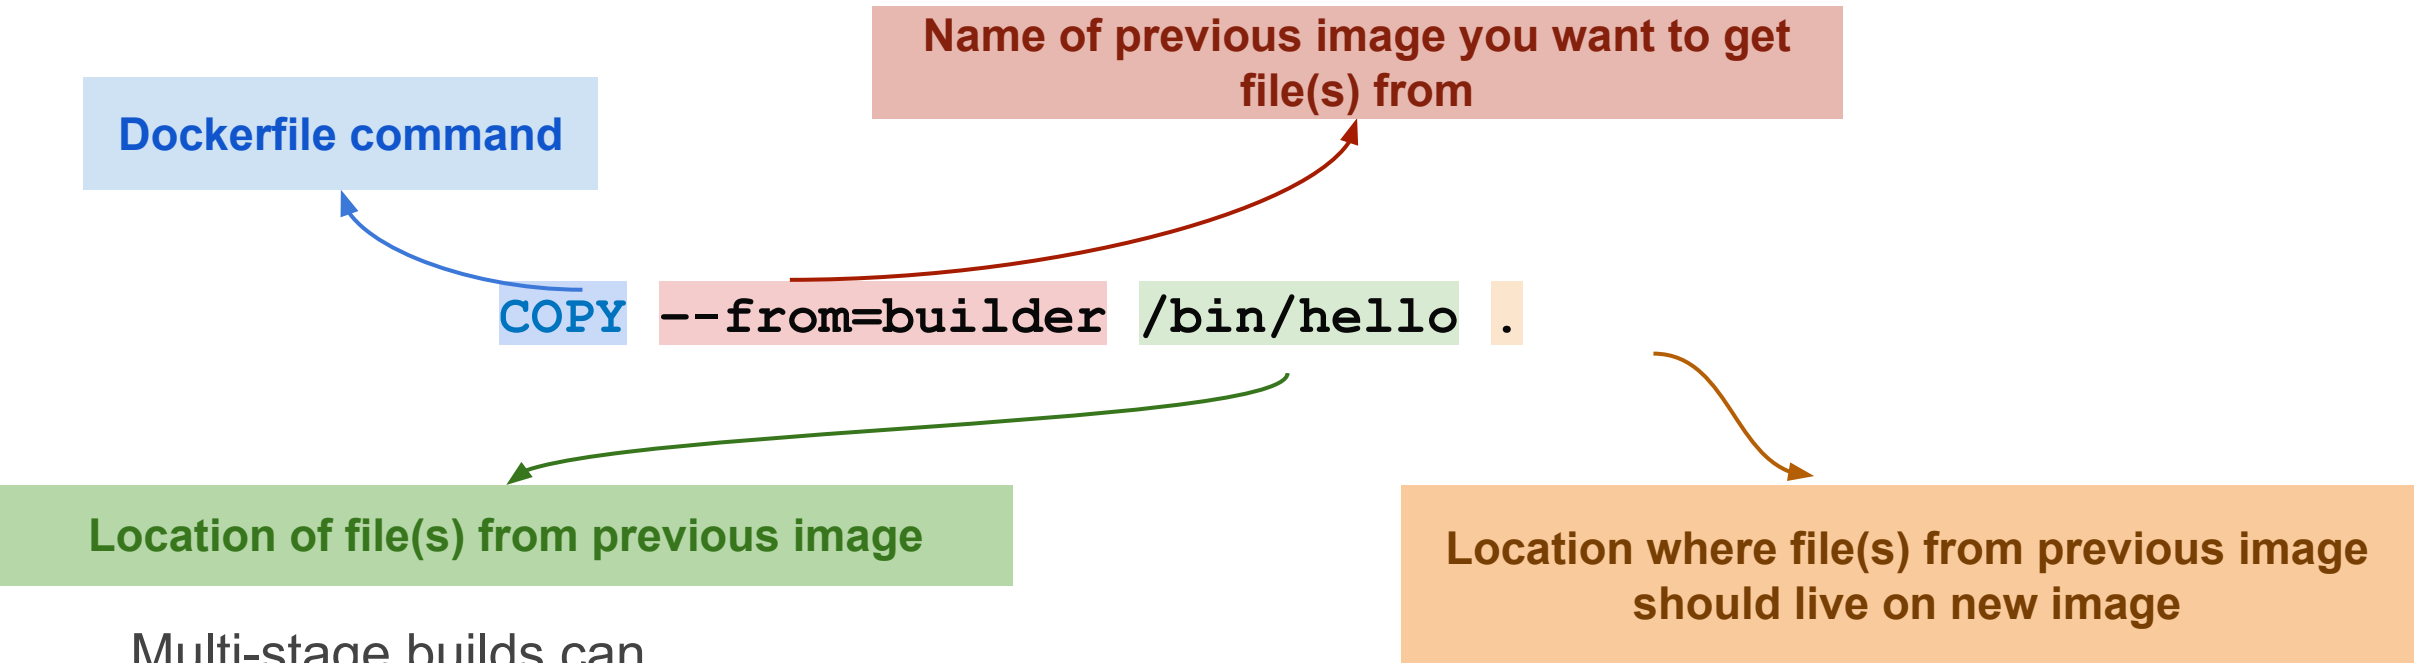
**Answer:** Usually no! Usually not! We typically only need the executable and its runtime dependencies.

Multi-stage builds allow us to bring **only what we need** into the final Docker image.

# Pro Tip 1: multi-stage builds

Multi-stage builds in Docker allow the use of multiple **FROM** statements in a single **Dockerfile**.

Each stage can bring in files from the previous stage using the **COPY** instruction.



Multi-stage builds can

- decrease container size
- improve security
- leverage different base images for each stage while preserving only the final image.

[multi-stage builds](#)

# Pro Tip 1: multi-stage builds: example

---

```
# First Stage: Building the Python application
```

```
FROM python:3.8 AS build-env
```

```
WORKDIR /app
```

```
COPY BLAH BLAH
```

```
RUN BLAH BLAH
```

```
# Second Stage: Copy the dependencies and run the application
```

```
FROM python:3.8-slim
```

```
COPY --from=build-env /usr/local/lib/python3.8/site-packages
```

```
/usr/local/lib/python3.8/site-packages
```

```
WORKDIR /app
```

```
COPY BLAH BLAH
```

```
CMD BLAH BLAH
```



## Pro Tip 2: multi-platform images

---

When building a more complicated Docker image, there is a small chance the specific platform (OS and CPU architecture) on your machine causes issues when sharing the Docker image with someone on a different machine

**Example:** Building a complex image on M1 Mac (linux/arm64) and trying to run the image on an older Macbook (linux/amd64)

Error message to look out for

The requested image's platform (linux/arm64) does not match the detected host platform (linux/amd64) and no specific platform was requested

**Solution:**

- Use the `--platform` flag within the FROM command in your Dockerfile to specify the target OS and CPU architecture for the build output

# Pro Tip 2: multi-platform images

---

Another way is to use buildx

```
# Check if buildx is available
```

```
docker buildx version
```

```
# Create a new builder instance named multi_platform.
```

```
docker buildx create --use --name multi_platform
```

```
# Create and load image for linux/arm64
```

```
docker buildx build --platform linux/arm64 -t print-platform-arm64 --load .
```

```
# Create and load image for linux/amd64
```

```
docker buildx build --platform linux/amd64 -t print-platform-amd64 --load .
```

```
# Run image container for linux/arm64
```

```
docker run --rm print-platform-arm64
```

```
# Run image container for linux/amd64
```

```
docker run --rm print-platform-amd64
```

## Pro Tip 3: **Docker Compose**

---

Docker Compose is a tool for defining and running multi-container Docker applications.

With Docker Compose, you can use a simple YAML file to configure your application's services, networks, and volumes, and then start everything with a single command.

We go through docker compose in more details next lecture.

## Logistics/Reminders

- If you have formed groups - please update [group info sheet](#)
- Please fill out survey -  
<https://canvas.harvard.edu/courses/136127/assignments/866239>  
(Survey responses have been updated)
- Office Hours details here -  
<https://edstem.org/us/courses/58478/discussion/5229430>





Now check: <https://formaggio.me/>